

In [442]:

```
%matplotlib inline
import pandas as pd
import numpy as np
from scipy.stats.stats import pearsonr
from pandas.tseries.holiday import USFederalHolidayCalendar
from pandas.tseries.offsets import CustomBusinessDay
from datetime import datetime
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
import seaborn as sns
```

In [235]:

```
trips = pd.read_csv('trip_data.csv')
stations = pd.read_csv('station_data.csv')
weather = pd.read_csv('weather_data.csv')
```

Data Exploration

Trip df & Station df

I'm going to summarize the data in trips df and station df and combine them into one df. To that end, I'll create some helper functions

In [124]:

```
def create_station_mapping(station_data):
    """
    Create a mapping from station IDs to cities, returning the
    result as a dictionary.
    """
    station_map = {}
    for data_file in station_data:
        with open(data_file, 'r') as f_in:

            weather_reader = csv.DictReader(f_in)

            for row in weather_reader:
                station_map[row['Id']] = row['Name']
    return station_map
```

In [372]:

```
def summarise_data(trip_in, station_data, trip_out):

    station_map = create_station_mapping(station_data)

    with open(trip_out, 'w') as f_out:
        out_colnames = ['start_date', 'start_year',
                        'start_month', 'start_hour', 'end_date', 'end_hour', 'weekday',
                        'start_station', 'end_station', 'subscription_type']

        trip_writer = csv.DictWriter(f_out, fieldnames = out_colnames)
        trip_writer.writeheader()

        for data_file in trip_in:
            with open(data_file, 'r') as f_in:
                trip_reader = csv.DictReader(f_in)

                # collect data from and process each row
                for row in trip_reader:
                    new_point = {}

                    trip_date = datetime.strptime(row['Start Date'], '%d/%m/%Y %
H:%M')

                    new_point['start_date'] = trip_date.strftime('%Y-%m-%d')
                    new_point['start_year'] = trip_date.strftime('%Y')
                    new_point['start_month'] = trip_date.strftime('%m')
                    new_point['start_hour'] = trip_date.strftime('%H')
                    new_point['weekday'] = trip_date.strftime('%a')

                    trip_enddate = datetime.strptime(row['End Date'], '%d/%m/%Y
%H:%M')

                    new_point['end_date'] = trip_enddate.strftime('%Y-%m-%d')
                    new_point['end_hour'] = trip_enddate.strftime('%H')

                    new_point['start_station'] = station_map[row['Start
Station']]

                    new_point['end_station'] = station_map[row['End Station']]
                    if 'Subscription Type' in row:
                        new_point['subscription_type'] = row['Subscription
Type']

                    else:
                        new_point['subscription_type'] = row['Subscriber Type']

                    trip_writer.writerow(new_point)
```

In [184]:

```
station_data = ['station_data.csv']
trip_in = ['trip_data.csv']
trip_out = 'trip_summary.csv'
summarise_data(trip_in, station_data, trip_out)
```

In [185]:

```
trips = pd.read_csv(trip_out)
```

In [186]:

```
# We check to see if there are any NA values and then see a sample of the first
few rows
trips.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 354152 entries, 0 to 354151
Data columns (total 10 columns):
start_date      354152 non-null object
start_year      354152 non-null int64
start_month     354152 non-null int64
start_hour      354152 non-null int64
end_date        354152 non-null object
end_hour        354152 non-null int64
weekday         354152 non-null object
start_station    354152 non-null object
end_station     354152 non-null object
subscription_type 354152 non-null object
dtypes: int64(4), object(6)
memory usage: 27.0+ MB
```

In [187]:

```
trips.head(5)
```

Out[187]:

	start_date	start_year	start_month	start_hour	end_date	end_hour	weekday	start_station
0	2015-08-31	2015	8	23	2015-08-31	23	Mon	Harry Bridge Plaza Building
1	2015-08-31	2015	8	23	2015-08-31	23	Mon	San Jose Shopping Center
2	2015-08-31	2015	8	23	2015-08-31	23	Mon	Post Kearney
3	2015-08-31	2015	8	23	2015-08-31	23	Mon	San Jose City Center
4	2015-08-31	2015	8	23	2015-08-31	23	Mon	Embassy at Fremont

This is the number of trips started at each station for a given date and hour

In [188]:

```
start_trips = trips.groupby(['start_station', 'start_date', 'start_hour'], as_index=False).count()
start_trips = start_trips.iloc[:, 0:4]
start_trips.rename(columns = {'start_year': 'start_trips'}, inplace = True)
start_trips.head()
```

Out[188]:

	start_station	start_date	start_hour	start_trips
0	2nd at Folsom	2014-09-01	14	1
1	2nd at Folsom	2014-09-01	16	1
2	2nd at Folsom	2014-09-02	6	1
3	2nd at Folsom	2014-09-02	7	1
4	2nd at Folsom	2014-09-02	8	2

This is the number of trips ended at each station for a given date and hour

In [197]:

```
end_trips = trips.groupby(['end_station', 'end_date', 'end_hour'], as_index=False).count()
end_trips = end_trips.iloc[:, 0:4]
end_trips.rename(columns = {'start_date': 'end_trips'}, inplace = True)
end_trips.head()
```

Out[197]:

	end_station	end_date	end_hour	end_trips
0	2nd at Folsom	2014-09-01	17	1
1	2nd at Folsom	2014-09-01	20	1
2	2nd at Folsom	2014-09-02	8	3
3	2nd at Folsom	2014-09-02	9	3
4	2nd at Folsom	2014-09-02	11	1

From these we can calculate the net rate of bike renting

In [405]:

```
test = pd.merge(start_trips, end_trips, how='outer', left_on=['start_station', 'start_date', 'start_hour'], right_on = ['end_station', 'end_date', 'end_hour'])
```

In [406]:

```
test.end_trips = test.end_trips.fillna(0)
```

In [407]:

```
test['net_trips'] = test.end_trips - test.start_trips
```

In [408]:

```
test = test.drop(['end_station', 'end_date', 'end_hour'], 1)
```

In [409]:

```
test = test.dropna()
```

In [410]:

```
test.head()
```

Out[410]:

	start_station	start_date	start_hour	start_trips	end_trips	net_trips
0	2nd at Folsom	2014-09-01	14.0	1.0	0.0	-1.0
1	2nd at Folsom	2014-09-01	16.0	1.0	0.0	-1.0
2	2nd at Folsom	2014-09-02	6.0	1.0	0.0	-1.0
3	2nd at Folsom	2014-09-02	7.0	1.0	0.0	-1.0
4	2nd at Folsom	2014-09-02	8.0	2.0	3.0	1.0

In [411]:

```
stns = pd.get_dummies(test.start_station)
```

In [412]:

```
test = test.merge(stns, left_index = True, right_index = True)
```

In [414]:

```
test = test.drop('start_station', 1)
```

Let's find out what are the top 10 most popular stations to start a ride

In [146]:

```
net_trips = start_trips.copy()
net_trips['end_trips'] = end_trips.end_trips
net_trips['net_trips'] = end_trips.end_trips - start_trips.start_trips
net_trips.rename(columns = {'start_station':'station'}, inplace = True)
net_trips.rename(columns = {'start_hour':'hour'}, inplace = True)
net_trips.head()
```

Out[146]:

	station	hour	start_trips	end_trips	net_trips
0	2nd at Folsom	0	17	15	-2.0
1	2nd at Folsom	1	9	1	-8.0
2	2nd at Folsom	2	5	1	-4.0
3	2nd at Folsom	4	70	3	-67.0
4	2nd at Folsom	5	125	212	87.0

In [381]:

```
total_trips = net_trips.groupby('station', as_index = False).sum()
total_trips.head()
top10_stations_start = list(total_trips.sort_values('start_trips',
ascending=0).station[:10])
top10_stations_start
```

Out[381]:

```
['San Francisco Caltrain (Townsend at 4th)',
'San Francisco Caltrain 2 (330 Townsend)',
'Harry Bridges Plaza (Ferry Building)',
'Temporary Transbay Terminal (Howard at Beale)',
'Embarcadero at Sansome',
'2nd at Townsend',
'Townsend at 7th',
'Steuart at Market',
'Market at 10th',
'Market at Sansome']
```

In []:

```
And top 10 stations to end a ride
```

In [382]:

```
top10_stations_end = list(total_trips.sort_values('end_trips', ascending=0).station[:10])
top10_stations_end
```

Out[382]:

```
['San Francisco Caltrain (Townsend at 4th)',
 'San Francisco Caltrain 2 (330 Townsend)',
 'Harry Bridges Plaza (Ferry Building)',
 '2nd at Townsend',
 'Embarcadero at Sansome',
 'Townsend at 7th',
 'Market at Sansome',
 'Temporary Transbay Terminal (Howard at Beale)',
 'Steuart at Market',
 'Powell Street BART']
```

Weather df

In [49]:

```
weather_df.head()
```

Out[49]:

	Date	Max TemperatureF	Mean TemperatureF	Min TemperatureF	Max Dew PointF	MeanDew PointF	Dewpo
0	01/09/2014	83.0	70.0	57.0	58.0	56.0	52.0
1	02/09/2014	72.0	66.0	60.0	58.0	57.0	55.0
2	03/09/2014	76.0	69.0	61.0	57.0	56.0	55.0
3	04/09/2014	74.0	68.0	61.0	57.0	57.0	56.0
4	05/09/2014	72.0	66.0	60.0	57.0	56.0	54.0

5 rows × 24 columns

In [271]:

```
weather.Date = pd.to_datetime(weather.Date, format='%d/%m/%Y')
```

In [273]:

```
weather.head()
```

Out[273]:

	Date	Max TemperatureF	Mean TemperatureF	Min TemperatureF	Max Dew PointF	MeanDew PointF	Min DewpointF
0	2014-09-01	83.0	70.0	57.0	58.0	56.0	52.0
1	2014-09-02	72.0	66.0	60.0	58.0	57.0	55.0
2	2014-09-03	76.0	69.0	61.0	57.0	56.0	55.0
3	2014-09-04	74.0	68.0	61.0	57.0	57.0	56.0
4	2014-09-05	72.0	66.0	60.0	57.0	56.0	54.0

5 rows × 24 columns

In [274]:

```
# Looks like we have multiple zip codes per date  
weather.Zip.unique()
```

Out[274]:

```
array([94107, 94063, 94301, 94041, 95113])
```


In [275]:

```
# See which zip code is "cleanest"  
for z in weather.Zip.unique():  
    print("Zip code:", z)  
    print(weather[weather.Zip == z].isnull().sum())  
    print()
```

```

Zip code: 94107
Date                                0
Max TemperatureF                    0
Mean TemperatureF                    0
Min TemperatureF                     0
Max Dew PointF                       0
MeanDew PointF                       0
Min DewpointF                        0
Max Humidity                         0
Mean Humidity                        0
Min Humidity                         0
Max Sea Level PressureIn             0
Mean Sea Level PressureIn            0
Min Sea Level PressureIn             0
Max VisibilityMiles                  0
Mean VisibilityMiles                 0
Min VisibilityMiles                  0
Max Wind SpeedMPH                    0
Mean Wind SpeedMPH                   0
Max Gust SpeedMPH                    6
PrecipitationIn                      0
CloudCover                           0
Events                              273
WindDirDegrees                       0
Zip                                  0
dtype: int64

```

```

Zip code: 94063
Date                                0
Max TemperatureF                    0
Mean TemperatureF                    0
Min TemperatureF                     0
Max Dew PointF                       0
MeanDew PointF                       0
Min DewpointF                        0
Max Humidity                         0
Mean Humidity                        0
Min Humidity                         0
Max Sea Level PressureIn             0
Mean Sea Level PressureIn            0
Min Sea Level PressureIn             0
Max VisibilityMiles                  0
Mean VisibilityMiles                 0
Min VisibilityMiles                  0
Max Wind SpeedMPH                    0
Mean Wind SpeedMPH                   0
Max Gust SpeedMPH                    209
PrecipitationIn                      0
CloudCover                           0
Events                              317
WindDirDegrees                       0
Zip                                  0
dtype: int64

```

```

Zip code: 94301
Date                                0
Max TemperatureF                     1
Mean TemperatureF                     1
Min TemperatureF                      1
Max Dew PointF                        47
MeanDew PointF                        47

```

Min DewpointF	47
Max Humidity	47
Mean Humidity	47
Min Humidity	47
Max Sea Level PressureIn	1
Mean Sea Level PressureIn	1
Min Sea Level PressureIn	1
Max VisibilityMiles	5
Mean VisibilityMiles	5
Min VisibilityMiles	5
Max Wind SpeedMPH	1
Mean Wind SpeedMPH	1
Max Gust SpeedMPH	313
PrecipitationIn	1
CloudCover	1
Events	324
WindDirDegrees	1
Zip	0
dtype: int64	

Zip code: 94041	
Date	0
Max TemperatureF	3
Mean TemperatureF	3
Min TemperatureF	3
Max Dew PointF	3
MeanDew PointF	3
Min DewpointF	3
Max Humidity	3
Mean Humidity	3
Min Humidity	3
Max Sea Level PressureIn	0
Mean Sea Level PressureIn	0
Min Sea Level PressureIn	0
Max VisibilityMiles	0
Mean VisibilityMiles	0
Min VisibilityMiles	0
Max Wind SpeedMPH	0
Mean Wind SpeedMPH	0
Max Gust SpeedMPH	7
PrecipitationIn	0
CloudCover	0
Events	311
WindDirDegrees	0
Zip	0
dtype: int64	

Zip code: 95113	
Date	0
Max TemperatureF	0
Mean TemperatureF	0
Min TemperatureF	0
Max Dew PointF	0
MeanDew PointF	0
Min DewpointF	0
Max Humidity	0
Mean Humidity	0
Min Humidity	0
Max Sea Level PressureIn	0
Mean Sea Level PressureIn	0
Min Sea Level PressureIn	0

```
Max VisibilityMiles      0
Mean VisibilityMiles     0
Min VisibilityMiles      0
Max Wind SpeedMPH        0
Mean Wind SpeedMPH       0
Max Gust SpeedMPH        6
PrecipitationIn          0
CloudCover               0
Events                  313
WindDirDegrees           0
Zip                      0
dtype: int64
```

In [276]:

```
# 94107 seems to be the best so I'll stick to this zip code from now on
weather = weather[weather.Zip == 94107]
```

In [277]:

```
weather.Events.unique()
```

Out[277]:

```
array([nan, 'Rain', 'Fog', 'Fog-Rain', 'Rain-Thunderstorm'], dtype=object)
```

In [278]:

```
# nan most likely means normal weather condition
weather.loc[weather.Events == 'Rain', 'Events'] = "Rain"
weather.loc[weather.Events.isnull(), 'Events'] = "Normal"
```

```
/Users/VAN/anaconda/lib/python3.6/site-packages/pandas/core/indexing.py:517: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy
self.obj[item] = s
```

In [279]:

```
weather.Events
```

Out[279]:

0	Normal
1	Normal
2	Normal
3	Normal
4	Normal
5	Normal
6	Normal
7	Normal
8	Normal
9	Normal
10	Normal
11	Normal
12	Normal
13	Normal
14	Normal
15	Normal
16	Rain
17	Rain
18	Normal
19	Normal
20	Normal
21	Normal
22	Rain
23	Rain
24	Rain
25	Normal
26	Normal
27	Normal
28	Normal
29	Normal
...	
335	Normal
336	Normal
337	Normal
338	Normal
339	Rain-Thunderstorm
340	Normal
341	Normal
342	Normal
343	Normal
344	Normal
345	Normal
346	Normal
347	Normal
348	Normal
349	Normal
350	Normal
351	Normal
352	Normal
353	Normal
354	Normal
355	Normal
356	Normal
357	Normal
358	Normal
359	Normal
360	Normal
361	Normal
362	Rain
363	Normal

365

Normal

364 Normal

Name: Events, Length: 365, dtype: object

In [280]:

```
# Convert values to new features
events = pd.get_dummies(weather.Events)
weather = weather.merge(events, left_index = True, right_index = True)
```

In [283]:

```
# Remove features we no longer need
weather = weather.drop(['Events', 'Zip'], 1)
```

In [284]:

```
weather.head()
```

Out[284]:

	Date	Max TemperatureF	Mean TemperatureF	Min TemperatureF	Max Dew PointF	MeanDew PointF	Min DewpointF
0	2014-09-01	83.0	70.0	57.0	58.0	56.0	52.0
1	2014-09-02	72.0	66.0	60.0	58.0	57.0	55.0
2	2014-09-03	76.0	69.0	61.0	57.0	56.0	55.0
3	2014-09-04	74.0	68.0	61.0	57.0	57.0	56.0
4	2014-09-05	72.0	66.0	60.0	57.0	56.0	54.0

5 rows × 27 columns

In [285]:

```
weather.isnull().sum()
```

Out[285]:

```
Date                                0
Max TemperatureF                    0
Mean TemperatureF                    0
Min TemperatureF                     0
Max Dew PointF                       0
MeanDew PointF                       0
Min DewpointF                        0
Max Humidity                         0
Mean Humidity                        0
Min Humidity                         0
Max Sea Level PressureIn              0
Mean Sea Level PressureIn             0
Min Sea Level PressureIn              0
Max VisibilityMiles                   0
Mean VisibilityMiles                  0
Min VisibilityMiles                   0
Max Wind SpeedMPH                     0
Mean Wind SpeedMPH                    0
Max Gust SpeedMPH                     6
PrecipitationIn                       0
CloudCover                           0
WindDirDegrees                        0
Fog                                   0
Fog-Rain                             0
Normal                               0
Rain                                  0
Rain-Thunderstorm                     0
dtype: int64
```

In [286]:

```
# Still have NAs in max gust speed. However, max wind speed and max gust speed i
s likely to be
# correlated => can use the former to fill the latter. Verify correlation with P
earson's corr
```

```
print(pearsonr(weather['Max Wind SpeedMPH'][weather['Max Gust SpeedMPH'] >= 0],
              weather['Max Gust SpeedMPH'][weather['Max Gust SpeedMPH'] >= 0]))
```

```
(0.75997233822788968, 8.7711694198354035e-69)
```

In [287]:

```
#For each value of max_wind, find the median max_gust
weather.loc[weather['Max Gust SpeedMPH'].isnull(), 'Max Gust SpeedMPH'] = weathe
r.groupby('Max Wind SpeedMPH')['Max Gust SpeedMPH'].apply(lambda x: x.fillna(x.m
edian()))
```


In [288]:

```
weather.dtypes
```

Out[288]:

Date	datetime64[ns]
Max TemperatureF	float64
Mean TemperatureF	float64
Min TemperatureF	float64
Max Dew PointF	float64
MeanDew PointF	float64
Min DewpointF	float64
Max Humidity	float64
Mean Humidity	float64
Min Humidity	float64
Max Sea Level PressureIn	float64
Mean Sea Level PressureIn	float64
Min Sea Level PressureIn	float64
Max VisibilityMiles	float64
Mean VisibilityMiles	float64
Min VisibilityMiles	float64
Max Wind SpeedMPH	float64
Mean Wind SpeedMPH	float64
Max Gust SpeedMPH	float64
PrecipitationIn	float64
CloudCover	float64
WindDirDegrees	float64
Fog	uint8
Fog-Rain	uint8
Normal	uint8
Rain	uint8
Rain-Thunderstorm	uint8
dtype:	object

In [289]:

```
weather.head()
```

Out[289]:

	Date	Max TemperatureF	Mean TemperatureF	Min TemperatureF	Max Dew PointF	MeanDew PointF	Min DewpointF
0	2014-09-01	83.0	70.0	57.0	58.0	56.0	52.0
1	2014-09-02	72.0	66.0	60.0	58.0	57.0	55.0
2	2014-09-03	76.0	69.0	61.0	57.0	56.0	55.0
3	2014-09-04	74.0	68.0	61.0	57.0	57.0	56.0
4	2014-09-05	72.0	66.0	60.0	57.0	56.0	54.0

5 rows × 27 columns

In [290]:

```
num_trips.head()
```

Out[290]:

	Date	num_trips
0	2014-09-01	368
1	2014-09-02	1319
2	2014-09-03	1404
3	2014-09-04	1389
4	2014-09-05	1265

In [291]:

```
weather.head()
```

Out[291]:

	Date	Max TemperatureF	Mean TemperatureF	Min TemperatureF	Max Dew PointF	MeanDew PointF	Min DewpointF
0	2014-09-01	83.0	70.0	57.0	58.0	56.0	52.0
1	2014-09-02	72.0	66.0	60.0	58.0	57.0	55.0
2	2014-09-03	76.0	69.0	61.0	57.0	56.0	55.0
3	2014-09-04	74.0	68.0	61.0	57.0	57.0	56.0
4	2014-09-05	72.0	66.0	60.0	57.0	56.0	54.0

5 rows × 27 columns

In [415]:

```
test['Date'] = test['start_date']
test['num_trips'] = test['net_trips']
```

In [416]:

```
test = test.groupby('Date', as_index=False).sum()
```

In [417]:

```
test = test.drop(['start_hour', 'start_trips', 'net_trips'], 1)
```

In [418]:

```
test = test.drop(['end_trips'], 1)
```

In [419]:

```
test.head()
```

Out[419]:

	Date	2nd at Folsom	2nd at South Park	2nd at Townsend	5th at Howard	Adobe on Almaden	Arena Green / SAP Center	Beale at Market	Broadway St at Battery St	Br
0	2014-09-01	2	4	4	2	0	1	3	4	0
1	2014-09-02	15	12	14	11	1	2	15	12	0
2	2014-09-03	16	13	12	9	1	3	12	9	1
3	2014-09-04	10	13	16	7	2	2	14	9	0
4	2014-09-05	14	11	13	10	2	0	14	10	0

5 rows × 73 columns

In [420]:

```
train = test.merge(weather, on = test.Date)
```

In [421]:

```
train.head()
```

Out[421]:

	Date_x	2nd at Folsom	2nd at South Park	2nd at Townsend	5th at Howard	Adobe on Almaden	Arena Green / SAP Center	Beale at Market	Broadway St at Battery St	
0	2014-09-01	2	4	4	2	0	1	3	4	C
1	2014-09-02	15	12	14	11	1	2	15	12	C
2	2014-09-03	16	13	12	9	1	3	12	9	1
3	2014-09-04	10	13	16	7	2	2	14	9	C
4	2014-09-05	14	11	13	10	2	0	14	10	C

5 rows × 100 columns

In [422]:

```
train['date'] = train['Date_x']
train.drop(['Date_y', 'Date_x'], 1, inplace= True)
```

In [423]:

```
train.head()
```

Out[423]:

	2nd at Folsom	2nd at South Park	2nd at Townsend	5th at Howard	Adobe on Almaden	Arena Green / SAP Center	Beale at Market	Broadway St at Battery St	Broadway at Main
0	2	4	4	2	0	1	3	4	0
1	15	12	14	11	1	2	15	12	0
2	16	13	12	9	1	3	12	9	1
3	10	13	16	7	2	2	14	9	0
4	14	11	13	10	2	0	14	10	0

5 rows × 99 columns

Add Special Date Features

In [424]:

```
#Find all of the holidays during our time span
calendar = USFederalHolidayCalendar()
holidays = calendar.holidays(start=train.date.min(), end=train.date.max())
```

In [425]:

```
holidays
```

Out[425]:

```
DatetimeIndex(['2014-09-01', '2014-10-13', '2014-11-11', '2014-11-27',
               '2014-12-25', '2015-01-01', '2015-01-19', '2015-02-16',
               '2015-05-25', '2015-07-03'],
              dtype='datetime64[ns]', freq=None)
```

In [426]:

```
#Find all of the business days in our time span
us_bd = CustomBusinessDay(calendar=USFederalHolidayCalendar())
business_days = pd.DatetimeIndex(start=train.date.min(), end=train.date.max(), freq=us_bd)
```

In [427]:

business_days

Out[427]:

```
DatetimeIndex(['2014-09-02', '2014-09-03', '2014-09-04', '2014-09-05',
               '2014-09-08', '2014-09-09', '2014-09-10', '2014-09-11',
               '2014-09-12', '2014-09-15',
               ...,
               '2015-08-18', '2015-08-19', '2015-08-20', '2015-08-21',
               '2015-08-24', '2015-08-25', '2015-08-26', '2015-08-27',
               '2015-08-28', '2015-08-31'],
              dtype='datetime64[ns]', length=251, freq='C')
```

In [428]:

```
business_days = pd.to_datetime(business_days, format='%Y/%m/%d').date
holidays = pd.to_datetime(holidays, format='%Y/%m/%d').date
```

In [429]:

```
train['business_day'] = train.date.isin(business_days)
train['holiday'] = train.date.isin(holidays)
```

In [430]:

train.head()

Out[430]:

	2nd at Folsom	2nd at South Park	2nd at Townsend	5th at Howard	Adobe on Almaden	Arena Green / SAP Center	Beale at Market	Broadway St at Battery St	Broadway at Mai
0	2	4	4	2	0	1	3	4	0
1	15	12	14	11	1	2	15	12	0
2	16	13	12	9	1	3	12	9	1
3	10	13	16	7	2	2	14	9	0
4	14	11	13	10	2	0	14	10	0

5 rows × 101 columns

In [431]:

```
#Convert True to 1 and False to 0
train.business_day = train.business_day.map(lambda x: 1 if x == True else 0)
train.holiday = train.holiday.map(lambda x: 1 if x == True else 0)
```

In [432]:

```
train['year'] = pd.to_datetime(train['date']).dt.year
train['month'] = pd.to_datetime(train['date']).dt.month
train['weekday'] = pd.to_datetime(train['date']).dt.weekday
```

In [433]:

```
labels = train.num_trips
train = train.drop(['num_trips', 'date'], 1)
```

Modeling approach

First split the data into training and test set

In [435]:

```
X_train, X_test, y_train, y_test = train_test_split(train, labels,
test_size=0.2, random_state = 2)
```

Linear Regression

In [444]:

```
from sklearn.linear_model import LinearRegression,Ridge,Lasso
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import mean_squared_error
import warnings
pd.options.mode.chained_assignment = None
warnings.filterwarnings("ignore", category=DeprecationWarning)

# Initialize logistic regression model
lModel = LinearRegression()

# Train the model
lModel.fit(X = X_train,y = y_train)

# Make predictions
preds = lModel.predict(X= X_train)

print("RMSE Value For Linear Regression: ", mean_squared_error(y_train,
preds)**0.5)
```

RMSE Value For Linear Regression: 16.5663679122

Regularization Model - Ridge

In [447]:

```
def rmse(y, y_):
    return mean_squared_error(y,y_)**0.5

ridge_m_ = Ridge()
ridge_params_ = { 'max_iter':[3000], 'alpha':[0.1, 1, 2, 3, 4, 10,
30,100,200,300,400,800,900,1000]}
rmse_scorer = metrics.make_scorer(rmse, greater_is_better=False)
grid_ridge_m = GridSearchCV( ridge_m_,
                             ridge_params_,
                             scoring = rmse_scorer,
                             cv=5)


grid_ridge_m.fit( X_train, y_train )
preds = grid_ridge_m.predict(X= X_train)
print (grid_ridge_m.best_params_)
print ("RMSE Value For Ridge Regression: ",rmse(y_train,preds))

fig,ax= plt.subplots()
fig.set_size_inches(12,5)
df = pd.DataFrame(grid_ridge_m.grid_scores_)
df["alpha"] = df["parameters"].apply(lambda x:x["alpha"])
df["rmse"] = df["mean_validation_score"].apply(lambda x:-x)
sns.pointplot(data=df,x="alpha",y="rmse",ax=ax)
```

```
{'alpha': 800, 'max_iter': 3000}
RMSE Value For Ridge Regression:  19.3441269662
```

Out[447]:

<matplotlib.axes._subplots.AxesSubplot at 0x121d1ad68>



Regularization Model - Lasso

In [450]:

```
lasso_m_ = Lasso()

alpha = 1/np.array([0.1, 1, 2, 3, 4, 10, 30,100,200,300,400,800,900,1000])
lasso_params_ = { 'max_iter':[3000],'alpha':alpha}

grid_lasso_m = GridSearchCV( lasso_m_,lasso_params_,scoring = rmse_scorer,cv=5)
grid_lasso_m.fit( X_train, y_train )
preds = grid_lasso_m.predict(X= X_train)
print (grid_lasso_m.best_params_)
print ("RMSE Value For Lasso Regression: ",rmse(y_train,preds))

fig,ax= plt.subplots()
fig.set_size_inches(12,5)
df = pd.DataFrame(grid_lasso_m.grid_scores_)
df["alpha"] = df["parameters"].apply(lambda x:x["alpha"])
df["rmse"] = df["mean_validation_score"].apply(lambda x:-x)
sns.pointplot(data=df,x="alpha",y="rmse",ax=ax)
```

```
/Users/VAN/anaconda/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:484: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Fitting data with very small alpha may cause precision problems.  
ConvergenceWarning)  
/Users/VAN/anaconda/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:484: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Fitting data with very small alpha may cause precision problems.  
ConvergenceWarning)  
/Users/VAN/anaconda/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:484: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Fitting data with very small alpha may cause precision problems.  
ConvergenceWarning)  
/Users/VAN/anaconda/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:484: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Fitting data with very small alpha may cause precision problems.  
ConvergenceWarning)  
/Users/VAN/anaconda/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:484: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Fitting data with very small alpha may cause precision problems.  
ConvergenceWarning)  
/Users/VAN/anaconda/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:484: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Fitting data with very small alpha may cause precision problems.  
ConvergenceWarning)  
/Users/VAN/anaconda/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:484: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Fitting data with very small alpha may cause precision problems.  
ConvergenceWarning)  
/Users/VAN/anaconda/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:484: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Fitting data with very small alpha may cause precision problems.  
ConvergenceWarning)  
/Users/VAN/anaconda/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:484: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Fitting data with very small alpha may cause precision problems.  
ConvergenceWarning)
```

```
l/coordinate_descent.py:484: ConvergenceWarning: Objective did not c
onverge. You might want to increase the number of iterations. Fittin
g data with very small alpha may cause precision problems.
```

```
ConvergenceWarning)
```

```
/Users/VAN/anaconda/lib/python3.6/site-packages/sklearn/linear_mode
l/coordinate_descent.py:484: ConvergenceWarning: Objective did not c
onverge. You might want to increase the number of iterations. Fittin
g data with very small alpha may cause precision problems.
```

```
ConvergenceWarning)
```

```
/Users/VAN/anaconda/lib/python3.6/site-packages/sklearn/linear_mode
l/coordinate_descent.py:484: ConvergenceWarning: Objective did not c
onverge. You might want to increase the number of iterations. Fittin
g data with very small alpha may cause precision problems.
```

```
ConvergenceWarning)
```

```
/Users/VAN/anaconda/lib/python3.6/site-packages/sklearn/linear_mode
l/coordinate_descent.py:484: ConvergenceWarning: Objective did not c
onverge. You might want to increase the number of iterations. Fittin
g data with very small alpha may cause precision problems.
```

```
ConvergenceWarning)
```

```
/Users/VAN/anaconda/lib/python3.6/site-packages/sklearn/linear_mode
l/coordinate_descent.py:484: ConvergenceWarning: Objective did not c
onverge. You might want to increase the number of iterations. Fittin
g data with very small alpha may cause precision problems.
```

```
ConvergenceWarning)
```

```
/Users/VAN/anaconda/lib/python3.6/site-packages/sklearn/linear_mode
l/coordinate_descent.py:484: ConvergenceWarning: Objective did not c
onverge. You might want to increase the number of iterations. Fittin
g data with very small alpha may cause precision problems.
```

```
ConvergenceWarning)
```

```
/Users/VAN/anaconda/lib/python3.6/site-packages/sklearn/linear_mode
l/coordinate_descent.py:484: ConvergenceWarning: Objective did not c
onverge. You might want to increase the number of iterations. Fittin
g data with very small alpha may cause precision problems.
```

```
ConvergenceWarning)
```

```
/Users/VAN/anaconda/lib/python3.6/site-packages/sklearn/linear_mode
l/coordinate_descent.py:484: ConvergenceWarning: Objective did not c
onverge. You might want to increase the number of iterations. Fittin
g data with very small alpha may cause precision problems.
```

```
ConvergenceWarning)
```

```
/Users/VAN/anaconda/lib/python3.6/site-packages/sklearn/linear_mode
l/coordinate_descent.py:484: ConvergenceWarning: Objective did not c
onverge. You might want to increase the number of iterations. Fittin
g data with very small alpha may cause precision problems.
```

```
ConvergenceWarning)
```

```
/Users/VAN/anaconda/lib/python3.6/site-packages/sklearn/linear_mode
l/coordinate_descent.py:484: ConvergenceWarning: Objective did not c
onverge. You might want to increase the number of iterations. Fittin
g data with very small alpha may cause precision problems.
```

```
ConvergenceWarning)
```

```
/Users/VAN/anaconda/lib/python3.6/site-packages/sklearn/linear_mode
l/coordinate_descent.py:484: ConvergenceWarning: Objective did not c
onverge. You might want to increase the number of iterations. Fittin
g data with very small alpha may cause precision problems.
```

```
ConvergenceWarning)
```

```
/Users/VAN/anaconda/lib/python3.6/site-packages/sklearn/linear_mode
l/coordinate_descent.py:484: ConvergenceWarning: Objective did not c
onverge. You might want to increase the number of iterations. Fittin
g data with very small alpha may cause precision problems.
```

```
ConvergenceWarning)
```

```
/Users/VAN/anaconda/lib/python3.6/site-packages/sklearn/linear_mode
l/coordinate_descent.py:484: ConvergenceWarning: Objective did not c
```

```
ConvergenceWarning)
/Users/VAN/anaconda/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:484: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Fitting
```